

# Service Level Objectives In Practice And At Scale

---

Carson Anderson

Weave



@carsonoid



@carson\_ops

# SLOs at Weave

# Simple, Pre-Defined Objectives

Deploy: production

CLEAR

SAVE

Request SLO Class **critical**

99.99 % of HTTP requests will complete without error

90 % of HTTP requests will complete in under 0.1 seconds

99 % of HTTP requests will complete in under 0.25 seconds

99.99 % of GRPC requests will complete without error

90 % of GRPC requests will complete in under 0.1 seconds

99 % of GRPC requests will complete in under 0.25 seconds

```
1  deploy:
2    production:
3      slo:
4        requestClass: critical
5
```

# Fully Customized Objectives

Deploy: production

CLEAR

SAVE

Request SLO Class *None*



*Inheriting from one or more default configurations*

*Includes customization so final class is **None-custom***

90 % of HTTP requests will complete without error

95 % of HTTP requests will complete in under 1 second

90 % of HTTP requests will complete in under 10 seconds

90 % of gRPC requests will complete without error

95 % of gRPC requests will complete in under 0.5 seconds

90 % of gRPC requests will complete in under 1 second

```
1 deploy:
2   production:
3     slo:
4       requestSpec:
5         grpc:
6           availability:
7             percent: "90"
8           latency:
9             tier1:
10              percent: "90"
11              seconds: "1"
12            tier2:
13              percent: "95"
14              seconds: "0.5"
15         http:
16           availability:
17             percent: "90"
18           latency:
19             tier1:
20              percent: "90"
21              seconds: "10"
22            tier2:
23              percent: "95"
24              seconds: "1"
25
```

# Generated Dashboards

Request SLOs for example-service - Class: None-custom

Links: [WGraph](#)

99% of gRPC requests will complete without error



99% of gRPC requests will complete in under 2.5 second(s)



90% of gRPC requests will complete in under 5 second(s)



# Generated Dashboards - Pt. 2

99% of HTTP requests will complete without error



99% of HTTP requests will complete in under 2.5 second(s)



90% of HTTP requests will complete in under 5 second(s)



# Generated Alerts

<input type="checkbox"/>	#96989	<div>x1</div>	<b>P3</b> [FIRING:1] [WARNING] example-service-http-latency-tier1-RequestSLOHighBudgetBurn in production	<div> Carson</div>	<div>OPEN</div>	<div>AckClose...</div>	Nov 6, 2021 3:21 PM (GMT-06:00)
<input type="checkbox"/>	#96914	<div>x1</div>	<b>P1</b> [FIRING:1] [CRITICAL] example-service-http-latency-tier1-RequestSLOCriticalBudgetBurn in production	<div> Carson</div>	<div>Carson Anderson</div>	<div>CLOSEDDelete</div>	Nov 5, 2021 11:36 PM (GMT-06:00)

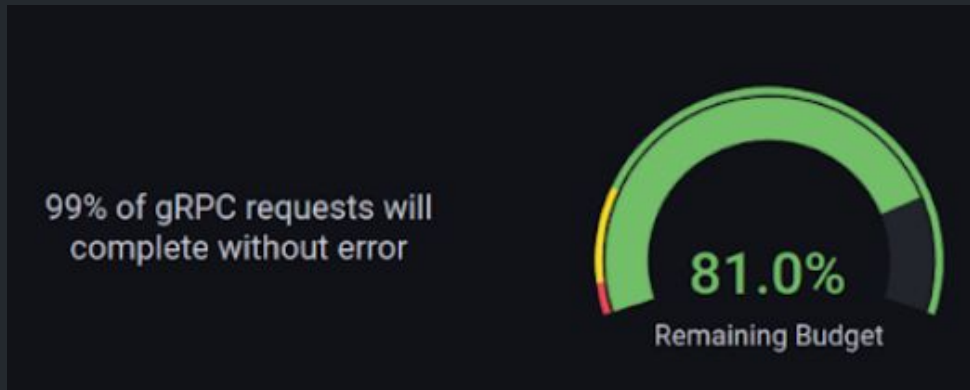
# So How Do I Get There?

- Calculating budgets
- Creating great alerts
- Handle scale
- Handle cost



# Calculating An SLO Budget

## The Math



# How to calculate an SLO

- Decide on an SLO window
  - Google SRE book: 30d
  - Weave: 28d (4 weeks)
- Calculate the number of objective "fails" allowed
  - Availability failure == response returns a server-side error
    - Ex: 500
  - Latency failure == response time < objective time
    - Ex: 2 seconds to respond
- Calculate the final budget % remaining
  - Based on total & budgeted

# Calculating Budgeted Errors

$$\text{Total} * .1 = \text{Budgeted}$$

1- (Objective/100)  
Ex: 1- (90/100)

$$50 * .1 = 5$$

## Calculating the SLO %

$$\frac{\text{Budgeted} - \text{Actual}}{\text{Budgeted}} = \text{Budget}$$

$$\frac{5 - 0}{5} = 100\%$$

# Basic SLO Budget Calculations

$$\text{Total} * .1 = \text{Budgeted}$$

1- (Objective/100)  
Ex: 1- (90/100)

$$\frac{\text{Budgeted} - \text{Actual}}{\text{Budgeted}} = \text{Budget}$$

+50 Good Requests

$$\text{Budget} = \frac{\text{Budgeted} - \text{Actual}}{\text{Budgeted}}$$

$$\text{Budgeted: } 5 = \text{Total: } 50 * .1$$

$$\frac{5 - 0}{5} = 100\%$$

+1 Bad Request

$$\text{Budget} = \frac{\text{Budgeted} - \text{Actual}}{\text{Budgeted}}$$

$$\text{Budgeted: } 5 = \text{Total: } 51 * .1$$

$$100\% = \frac{5 - 0}{5} \Rightarrow 80\% = \frac{5 - 1}{5}$$

+3 Bad Requests

$$\text{Budget} = \frac{\text{Budgeted} - \text{Actual}}{\text{Budgeted}}$$

$$\text{Budgeted: } 5 = \text{Total: } 54 * .1$$

$$80\% = \frac{5 - 1}{5} \Rightarrow 20\% = \frac{5 - 4}{5}$$



+1 Good Request

$$\text{Budget} = \frac{\text{Budgeted} - \text{Actual}}{\text{Budgeted}}$$

$$\text{Budgeted: } 5 = \text{Total: } 55 * .1$$

$$20\% = \frac{5 - 4}{5} \Rightarrow 20\% = \frac{5 - 4}{5}$$

+5 Good Requests

$$\text{Budget} = \frac{\text{Budgeted} - \text{Actual}}{\text{Budgeted}}$$

$$\text{Budgeted: } 6 = \text{Total: } 60 * .1$$

$$20\% = \frac{5 - 4}{5} \Rightarrow 33\% = \frac{6 - 4}{6}$$

Recover With Time

$$\text{Budget} = \frac{\text{Budgeted} - \text{Actual}}{\text{Budgeted}}$$

28d



$$25 * .1 = \text{Budgeted}:2$$

## Recover With Time

$$\text{Budget} = \frac{\text{Budgeted} - \text{Actual}}{\text{Budgeted}}$$

28d



$$\frac{2 - 4}{2} = -100\%$$

## Recover With Time

$$\text{Budget} = \frac{\text{Budgeted} - \text{Actual}}{\text{Budgeted}}$$

28d



$$-100\% = \frac{2 - 4}{2} \rightarrow 0\% = \frac{2 - 2}{2}$$

# Key Takeaways

- Budgets can burn FAST
- Two ways to recover:
  - Time
  - Handle more "Good" request, a lot of them

# Availability SLO Recording Rules Prometheus Specifics

# Prometheus Metric Crash-Course - counters

```
# HELP http_timer_count The total number of HTTP requests.
# TYPE http_timer_count counter
http_timer_count{app="api", code="200", method="GET"} 1
http_timer_count{app="api", code="200", method="POST"} 10
http_timer_count{app="api", code="200", method="PATCH"} 2
http_timer_count{app="frontend", code="200", method="GET"} 110
http_timer_count{app="frontend", code="404", method="GET"} 23
http_timer_count{app="frontend", code="500", method="GET"} 2
```



# Example Latency Histogram in Prometheus

```
# HELP http_timer_bucket A histogram of request duration
# TYPE http_timer_bucket histogram
http_timer_bucket{app="api", code="200", method="GET", le="+Inf"}
...
```

# Example Latency Histogram in Prometheus

## New Requests

- 1s latency

## Bucket Increments

```
le: .1      =  
le: .25     =  
le: 1       = X  
le: 5       = X  
le: +Inf    = X
```

# Example Latency Histogram in Prometheus

## New Requests

- 1s latency
- 5s latency

## Bucket Increments

le: .1 =  
le: .25 =  
le: 1 = X  
le: 5 = XX  
le: +Inf = XX

# Example Latency Histogram in Prometheus

## New Requests

- 1s latency
- 5s latency
- .1s latency

## Bucket Increments

le: .1 = X

le: .25 = X

le: 1 = XX

le: 5 = XXX

le: +Inf = XXX

# Prometheus Counter Reminder

```
# HELP http_timer_count The total number of HTTP requests.
# TYPE http_timer_count counter
http_timer_count{app="api", code="200", method="GET"} 1
http_timer_count{app="api", code="200", method="POST"} 10
http_timer_count{app="api", code="200", method="PATCH"} 2
http_timer_count{app="frontend", code="200", method="GET"} 110
http_timer_count{app="frontend", code="404", method="GET"} 23
http_timer_count{app="frontend", code="500", method="GET"} 2
```

# Prometheus Metric Crash-Course - histograms

```
# HELP http_timer_bucket A histogram of request duration
# TYPE http_timer_bucket histogram
http_timer_bucket{app="api", code="200", method="GET", l e="+Inf"} 1
http_timer_bucket{app="api", code="200", method="GET", l e="5"} 1
http_timer_bucket{app="api", code="200", method="POST", l e="+Inf"} 10
http_timer_bucket{app="api", code="200", method="POST", l e="5"} 7
http_timer_bucket{app="api", code="200", method="POST", l e="1"} 5
http_timer_bucket{app="api", code="200", method="POST", l e=".25"} 5
http_timer_bucket{app="api", code="200", method="POST", l e=".1"} 5
http_timer_bucket{app="api", code="200", method="PATCH", l e="+Inf"} 2
http_timer_bucket{app="frontend", code="200", method="GET", l e="+Inf"} 110
http_timer_bucket{app="frontend", code="200", method="GET", l e="5"} 101
http_timer_bucket{app="frontend", code="200", method="GET", l e="1"} 100
http_timer_bucket{app="frontend", code="200", method="GET", l e=".25"} 97
http_timer_bucket{app="frontend", code="200", method="GET", l e=".1"} 96
http_timer_bucket{app="frontend", code="404", method="GET", l e="+Inf"} 23
...
```

# FYI there is some duplication

```
# sum all variants of the count
sum(http_timer_count{app="frontend"}) = 135
```

```
# sum all variants of the catch-all bucket
sum(http_timer_bucket{app="frontend", le="+Inf"}) = 135
```

```
# some tooling sets a dedicated "bucket_count" metric
http_timer_bucket_count{app="frontend"} = 135
```

# Which Metrics to use?

`grpc_timer_count` & `http_timer_count`

- total count of requests

`grpc_timer_bucket` & `http_timer_bucket`

- histogram of requests, bucketed by time

`grpc_timer_bucket_count` & `http_timer_bucket_count`

- total number of requests through histogram



# Key Takeaways

- Histograms in Prometheus are **cumulative**
  - ALL requests will always get counted in "+Inf"
- Weave tooling: ``http_timer_bucket`` has everything
  - No need for any ``*_count`` metrics
  - Can dramatically cut total # of series in

# grpc\_timer\_bucket{"+Inf"} == grpc\_timer\_count



# Availability SLO Recording Rules ... The Easy Way

# Instantaneous Error Budget Rule For All Apps

```
(  
  ( sum(http_timer_bucket{le="+Inf"}) by (app) * .1 )  
  -  
  (  
    ( sum(http_timer_bucket{le="+Inf"}) by (app) )  
    -  
    ( sum(http_timer_bucket{le="+Inf",code!~"5.."}) by (app) )  
  )  
)  
/  
( sum (http_timer_bucket{le="+Inf"}) by (app) * .1 )
```

# 28d Error Budget Rule For All Apps

```
(  
  ( sum(increase(http_timer_bucket{le="+Inf"}[28d])) by (app) * .1 )  
  -  
  (  
    ( sum(increase(http_timer_bucket{le="+Inf"}[28d])) by (app) )  
    -  
    ( sum(increase(http_timer_bucket{le="+Inf",code!~"5.."}[28d])) by (app) )  
  )  
)  
/  
( sum(increase(http_timer_bucket{le="+Inf"}[28d])) by (app) * .1 )
```

# Lots of Duplication, Compute, Confusion

```
(  
  ( sum(increase(http_timer_bucket{le="+Inf"}[28d])) by (app) * .1 )  
  -  
  (  
    ( sum(increase(http_timer_bucket{le="+Inf"}[28d])) by (app) )  
    -  
    ( sum(increase(http_timer_bucket{le="+Inf",code!~"5.."}[28d])) by (app) )  
  )  
)  
/  
( sum(increase(http_timer_bucket{le="+Inf"}[28d])) by (app) * .1 )
```

## ⚠ Important Prometheus Recording Behavior! ⚠

- Prometheus supports Rules
  - alert rules - used to create alerts (duh)
  - record rules - run queries and save results
- Rules are provided in groups
- Rules are evaluated periodically
- Groups are processed in a random order
- **Rules within a group are always processed in order**

# Example Availability Prometheus Rules

$$\text{Budget} = \frac{\text{Budgeted} - \text{Actual}}{\text{Budgeted}}$$

groups:

- name: slo:example:budget\_all\_in\_one  
rules:
  - record: slo:total  
expr: sum(increase(http\_timer\_bucket{le="+Inf"}[28d])) by (app)
  - record: slo:budgeted  
expr: slo:total \* .1
  - record: slo:occured  
expr: |  
slo:total -  
sum(increase(http\_timer\_bucket{le="+Inf",code!~"5.."}[28d])) by (app)
  - record: slo:budget  
expr: (slo:budgeted - slo:occured) / slo:budgeted



# Old rule for comparison

```
(
  ( sum(increase(http_timer_bucket{le="+Inf"}[28d])) by (app) * .1 )
-
  (
    ( sum(increase(http_timer_bucket{le="+Inf"}[28d])) by (app) )
    -
    ( sum(increase(http_timer_bucket{le="+Inf",code!~"5.."}[28d])) by (app) )
  )
)
/
( sum(increase(http_timer_bucket{le="+Inf"}[28d])) by (app) * .1 )
```

There are still problems!

# Problems With All-In-One Rules

- All Apps have to have the same SLO objective
- Calculations for long SLO periods can be too large for Prometheus to calculate
  - "error executing query: query processing would load too many samples into memory in query execution"

# Real, Auto-Generated Availability SLO Recording Rules

# Remember This?

Deploy: production

CLEAR

SAVE

Request SLO Class *None*

*Inheriting from one or more default configurations*

*Includes customization so final class is **None-custom***

90 % of HTTP requests will complete without error

95 % of HTTP requests will complete in under 1 second

90 % of HTTP requests will complete in under 10 seconds

90 % of GRPC requests will complete without error

95 % of GRPC requests will complete in under 0.5 seconds

90 % of GRPC requests will complete in under 1 second

```
1 deploy:
2   production:
3     slo:
4       requestSpec:
5         grpc:
6           availability:
7             percent: "90"
8           latency:
9             tier1:
10              percent: "90"
11              seconds: "1"
12            tier2:
13              percent: "95"
14              seconds: "0.5"
15         http:
16           availability:
17             percent: "90"
18           latency:
19             tier1:
20              percent: "90"
21              seconds: "10"
22            tier2:
23              percent: "95"
24              seconds: "1"
25
```

# Calculate the Total # of requests for the budget

```
- record: 'slo:example_service:request:http:availability:budget_total  
  expr: |  
    sum(increase(http_timer_bucket{app="example-service",le="+Inf"}[28d]))
```

# Calculate the total # of requests allowed to fail

```
- record: 'slo:example_service:request:http:availability:budget_allowed'  
  expr: |  
    (slo:example_service:request:http:availability:tier1:budget_total * .01)
```

Objective:  
"99% Must Not Fail"

# Calculate the total # of requests to actually fail

```
- record: 'slo:example_service:request:http:availability:budget_violations'  
expr: >  
  (slo:example_service:request:http:availability:tier1:budget_total -  
   sum(increase(  
     http_timer_bucket{app="example-service",code!~"5..",le="+Inf"}[28d]  
   ))  
  )
```



# Calculate Remaining Budget

$$\text{Budget} = \frac{\text{Budgeted} - \text{Actual}}{\text{Budgeted}}$$

```
# Calculate budget as: Budgeted - Actual / Budgeted
- record: 'slo:example_service:request:http:availability:error_budget'
  expr: >
    (
      slo:example_service:request:http:availability:budget_allowed -
      slo:example_service:request:http:availability:budget_violations
    ) /
    slo:example_service:request:http:availability:budget_allowed
```

# Final Rules - Availability

$$\text{Budget} = \frac{\text{Budgeted} - \text{Actual}}{\text{Budgeted}}$$

# Calculate the Total # of requests within a timeframe that count towards the budget

```
- record: 'slo:example_service:request:http:availability:budget_total'
  expr: |
    sum(increase(http_timer_bucket{app="example-service",le="+Inf"}[28d]))
```

# Calculate the total # of requests within a timeframe **Budgeted** to fail

```
- record: 'slo:example_service:request:http:availability:budget_allowed'
  expr: |
    (slo:example_service:request:http:availability:budget_total * .01)
```

# Calculate the total # of requests within a timeframe that **Actually** failed

```
- record: 'slo:example_service:request:http:availability:budget_violations'
  expr: >
    (slo:example_service:request:http:availability:budget_total -
    sum(increase(http_timer_bucket{app="example-service",code!~"5..",le="+Inf"}[28d])))
```

# Calculate budget as: **Budgeted - Actual / Budgeted**

```
- record: 'slo:example_service:request:http:availability:error_budget'
  expr: >
    (slo:example_service:request:http:availability:budget_allowed -
    slo:example_service:request:http:availability:budget_violations) /
    slo:example_service:request:http:availability:budget_allowed
```

# Real Latency SLO Recording Rules

# Final Rules - Latency

$$\text{Budget} = \frac{\text{Budgeted} - \text{Actual}}{\text{Budgeted}}$$

# Calculate the Total # of requests within a timeframe that count toward

```
- record: 'slo:example_service:request:http:latency:tier1:budget_total'
  expr: |
    sum(increase(http_timer_bucket{app="example-service",le="+Inf"}[28d]))
```

# Calculate the total # of requests within a timeframe **Budgeted** to fail

```
- record: 'slo:example_service:request:http:latency:tier1:budget_allowed'
  expr: |
    (slo:example_service:request:http:latency:tier1:budget_total * .01)
```

# Calculate the total # of requests within a timeframe that **Actually** failed

```
- record: 'slo:example_service:request:http:latency:tier1:budget_violations'
  expr: >
    (slo:example_service:request:http:latency:tier1:budget_total -
    sum(increase(http_timer_bucket{app="example-service",le="1"}[28d])))
```

# Calculate budget as: **Budgeted - Actual / Budgeted**

```
- record: 'slo:example_service:request:http:latency:tier1:error_budget' # <- The actual budget
  expr: >
    (slo:example_service:request:http:latency:tier1:budget_allowed -
    slo:example_service:request:http:latency:tier1:budget_violations) /
    slo:example_service:request:http:latency:tier1:budget_allowed
```

Change rule names

Weave has two "tiers" of latency:

- tier 1 = stretch goal
- tier 2 = required

99% must complete...

...in 1 second or less


# Dashboard Queries Get To Be Simple!

`slo:example_service:request:http:availability:error_budget{}`

99% of HTTP requests will complete without error



99% of HTTP requests will complete in under 2.5 second(s)



`slo:example_service:request:http:latency:tier1:error_budget{}`

90% of HTTP requests will complete in under 5 second(s)



`slo:example_service:request:http:latency:tier2:error_budget{}`

# SLO Alerting

# Naive SLO Alert

```
- alert: example_service:LowHTTPAvailabilityBudget
  expr: |
    slo:example_service:request:http:availability:error_budget < .5
```

# Alert Periods






# Goal: Detect Fast Drops of High %



# Goal: Detect Slow Burn Over Time



# Great Alerts Use "Burn Rate"



## Calculate A Burn Rate

$$\frac{\text{error\_rate}}{\text{request\_rate}} = \text{burn\_rate}$$

# Calculate A Burn Rate In Prometheus

$$\frac{\text{request\_rate} - \text{good\_rate}}{\text{request\_rate}} = \text{burn\_rate}$$

# Availability Burn Rate Recording Rule 5m Example

```
- expr: |
    (
        sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[5m]))
        -
        sum(rate(http_timer_bucket{app="example-app",code!~"5..",le="+Inf"}[5m]))
    )
    /
    sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[5m]))
record: 'slo:example_app:request:http:availability:rate_5m'
```

# Availability Burn Rate Recording Rules

```
- expr: >
  (sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[5m])) -
   sum(rate(http_timer_bucket{app="example-app",code!~"5..",le="+Inf"}[5m])))
  / sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[5m]))
record: 'slo:example_app:request:http:availability:rate_5m'

- expr: >
  (sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[30m])) -
   sum(rate(http_timer_bucket{app="example-app",code!~"5..",le="+Inf"}[30m])))
  / sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[30m]))
record: 'slo:example_app:request:http:availability:rate_30m'

- expr: >
  (sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[1h])) -
   sum(rate(http_timer_bucket{app="example-app",code!~"5..",le="+Inf"}[1h])))
  / sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[1h]))
record: 'slo:example_app:request:http:availability:rate_1h'

- expr: >
  (sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[6h])) -
   sum(rate(http_timer_bucket{app="example-app",code!~"5..",le="+Inf"}[6h])))
  / sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[6h]))
record: 'slo:example_app:request:http:availability:rate_6h'

- expr: >
  (sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[3d])) -
   sum(rate(http_timer_bucket{app="example-app",code!~"5..",le="+Inf"}[3d])))
  / sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[3d]))
record: 'slo:example_app:request:http:availability:rate_3d'
```

# Latency Burn Rate Recording Rule 5m Example

```
- expr: |
  (
    sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[5m]))
    -
    sum(rate(http_timer_bucket{app="example-app",le="5"}[5m]))
  )
  /
  sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[5m]))
record: 'slo:example_app:request:http:latency:tier1:rate_5m'
```

Remove `code` and  
add a specific `le`

`le` **MUST** match a  
bucket reported in  
metrics



# Latency Burn Rate Recording Rules

```
- expr: >
  (sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[5m])) -
   sum(rate(http_timer_bucket{app="example-app",le="5"}[5m]))) /
   sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[5m]))
record: 'slo:example_app:request:http:latency:tier1:rate_5m'

- expr: >
  (sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[30m])) -
   sum(rate(http_timer_bucket{app="example-app",le="5"}[30m]))) /
   sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[30m]))
record: 'slo:example_app:request:http:latency:tier1:rate_30m'

- expr: >
  (sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[1h])) -
   sum(rate(http_timer_bucket{app="example-app",le="5"}[1h]))) /
   sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[1h]))
record: 'slo:example_app:request:http:latency:tier1:rate_1h'

- expr: >
  (sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[6h])) -
   sum(rate(http_timer_bucket{app="example-app",le="5"}[6h]))) /
   sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[6h]))
record: 'slo:example_app:request:http:latency:tier1:rate_6h'

- expr: >
  (sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[3d])) -
   sum(rate(http_timer_bucket{app="example-app",le="5"}[3d]))) /
   sum(rate(http_timer_bucket{app="example-app",le="+Inf"}[3d]))
record: 'slo:example_app:request:http:latency:tier1:rate_3d'
```

# Now we can build smarter alerts

- alert: CriticalBudgetBurn

labels:

severity: page

expr: |

(

(

slo:example\_app:request:http:latency:tier1:rate\_1h > (13.44 \* .01)

and

slo:example\_app:request:http:latency:tier1:rate\_5m > (13.44 \* .01)

)

or

(


slo:example\_app:request:http:latency:tier1:rate\_6h > (5.6 \* .01)

and


slo:example\_app:request:http:latency:tier1:rate\_30m > (5.6 \* .01)

)

)




# Real World Example




### Budget Remaining



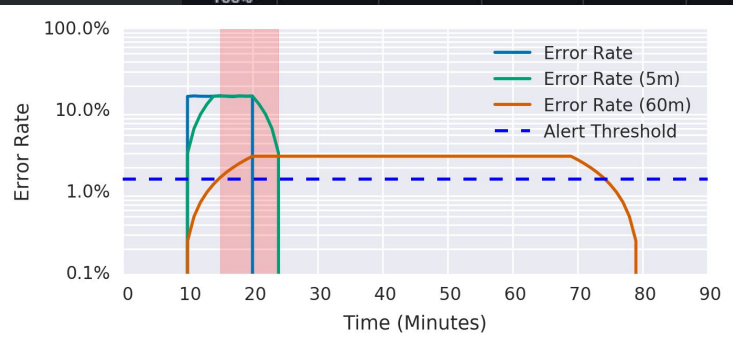

### Burn Rates



### Burn Rates



### Burn Rates w/Alert Overlay




# Alerting Breakdown


Budget Remaining



Burn Rates



Burn Rates w/Alert Overlay



# SLO Alerting Advice: Follow the SRE Book

- Use a 30d window
- Use the example rate windows (5m,30m,1h,6h,3d)
- Use the example alert rules

<https://sre.google/workbook/alerting-on-slos/>



Tone Shift!

Let's Talk K8s and Cloud Native

What about the Dashboards?

# Solution: Grafana Operator

- The [grafana-operator](#) has support for a dashboard CRD!

Unfortunately, it doesn't work work for grafana cloud

- We used the [operator-sdk](#) to build a custom operator at Weave which could
  - Take a sample dashboard JSON
  - Do some string substitutions
  - Apply to Grafana Cloud

## Request SLOs for \_\_app\_\_ - Class: \_\_class\_\_

[Links: WGraph](#)

99.9% of gRPC requests will complete without error






99% of gRPC requests will complete in under 5 second(s)



90% of gRPC requests will complete in under 1 second(s)



99.9% of HTTP requests will complete without error



# Example Dashboard YAML

```
apiVersion: custom.weave/v1beta1
kind: Dashboard
metadata:
  name: dashboard-sample
spec:
  grafana: local
  title: test-dash
  json: |
    {
      "title": "__app__ SLOs",
      "panels": [
        ...
      ],
      "timezone": "",
    }
  jsonReplacements:
    - find: __app__
      replace: example-app
```

# Generated Dashboards

Request SLOs for example-service - Class: None-custom

Links: [WGraph](#)

99% of gRPC requests will complete without error



99% of gRPC requests will complete in under 2.5 second(s)



90% of gRPC requests will complete in under 5 second(s)



# The Scale Issue

# Prometheus Scaling Concerns

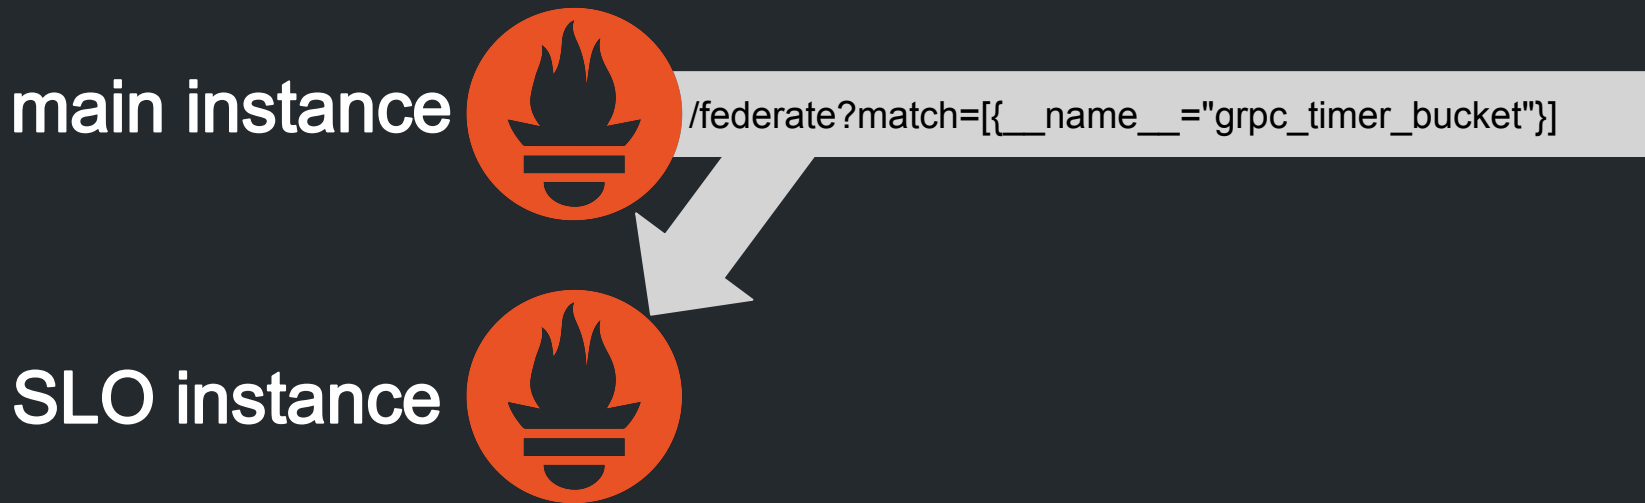
- The "main" Prometheus is huge and **busy**
- Our SLO is for 28 days but our main Prometheus is 14 days only
- SLO calculations can add load





# Solution: Federation!

- Prometheus exposes a "federate" endpoint
  - Takes zero or more match expressions to filter returned results
- Essentially: Scrape Prometheus metrics from Prometheus



# prometheus-operator to the rescue!

```
apiVersion: monitoring.coreos.com/v1
kind: Prometheus
metadata:
  name: prometheus-slo
  labels:
    prometheus: prometheus-slo
spec:
  replicas: 1
  retention: 28d
  serviceAccountName: prometheus-slo
  serviceMonitorSelector:
    matchLabels:
      scraper: prometheus-slo
```

# prometheus-operator to the rescue!

```
apiVersion: monitoring.coreos.com/v1
kind: PodMonitor
metadata:
  name: prometheus-slo-metrics
  labels:
    scraper: prometheus-slo
spec:
  jobLabel: prometheus-export-slo-rules
  namespaceSelector:
    matchNames:
      - prometheus
  selector:
    matchLabels:
      prometheus: main
  podMetricsEndpoints:
    - honorLabels: true
      interval: 30s
      params:
        'match[]':
          - '{__name__=~"grpc_timer_bucket"}'
          - '{__name__=~"http_timer_bucket"}'
      path: /federate
      targetPort: 9090
```

# The Retention Issue

# SLO Retention Concerns

- Management wants **2 years** of slo budget history
- Increasing even the slo Prometheus to 2y retention keeps **all** metrics forever

main instance



/federate?match=[{\_\_name\_\_="grpc\_timer\_bucket"}]

SLO instance



# Add Another Layer!

## main instance

- 14d retention
- all metrics for cluster
- no calculations required



/federate?match=[{\_\_name\_\_="grpc\_timer\_bucket"}]

## SLO instance

- 28d retention
- just slo metrics + calculations



/federate?match=[{\_\_name\_\_="slo:.\*"}]

## export instance

- 2y
- just slo recording rules



# prometheus-operator to the rescue!

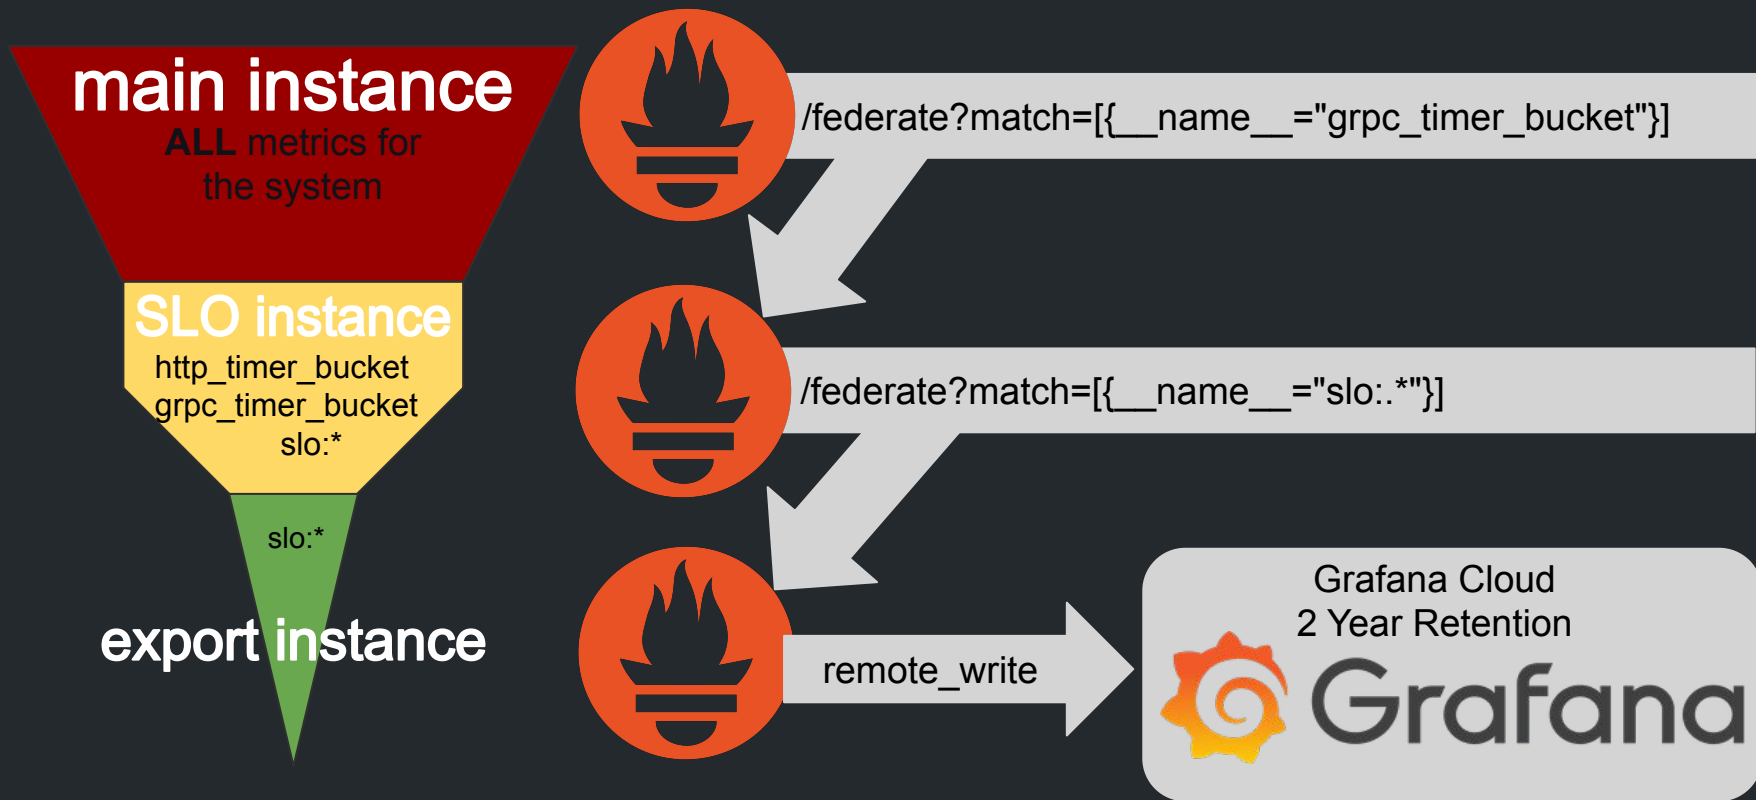
```
apiVersion: monitoring.coreos.com/v1
kind: Prometheus
metadata:
  name: prometheus-export
  labels:
    prometheus: prometheus-export
spec:
  replicas: 1
  retention: 2y
  serviceAccountName: prometheus-export
  serviceMonitorSelector:
    matchLabels:
      scraper: prometheus-export
```

# prometheus-operator to the rescue!

```
apiVersion: monitoring.coreos.com/v1
kind: PodMonitor
metadata:
  labels:
    scraper: prometheus-export
  name: prometheus-export-slo-rules
spec:
  jobLabel: prometheus-export-slo-rules
  selector:
    matchLabels:
      prometheus: slo
  namespaceSelector:
    matchNames:
      - prometheus-slo
  podMetricsEndpoints:
    - honorLabels: true
      interval: 30s
      params:
        'match[]':
          - '{__name__=~"slo:.*"}'
      path: /federate
      targetPort: 9090
```




# Prometheus Federation Series Count Funnel




# Grafana Cloud Billing Impact


# Ship ALL metrics




# Ship Recording Rules



# Ship only slo:\* rules



# Full Billing Timeline



# Final Advice

- Standardize your metrics
- Pick 1 or 2 objectives to start with
- Use a 30d window if you can
- Automate & Generate!

# Thank You!

## SLOs In Practice

---

Carson Anderson

Weave



@carsonoid



@carson\_ops